

AUTOMATED HATE SPEECH DETECTION IN ONLINE PLATFORMS



Group Members:
NIKITHA MATTAA
AKRATI SACHAN
MANIKANTA SEELAM

Agenda

01	Business Problem
02	Solution
03	Dataset Description
04	Data Visualization
05	Data Preprocessing
06	Tokenization and Embedding Techniques
07	Modeling
08	Evaluation metrics

Business Problem

Hate Speech Detection on Reddit

- Hate speech on Reddit can harm user safety, community health, and platform reputation.
- Effective detection is crucial for maintaining a safe and inclusive online environment.



Solution:



Building the Solution

Data Preparation:

- Collected and preprocessed Reddit comments.
- Addressed class imbalances using SMOTE

Model Development:

- Implemented ML models (SVM, SGD Classifier, Random forest etc.) and DL models (CNN, LSTM, BERT).
- BERT model achieved the best performance with 89% accuracy and an F1-score of 78%



Business Benefits

Enhanced User Safety:

- Protects users from harmful content, creating a safer environment.

Improved Community Health:

- Reduces toxicity, encouraging positive and constructive interactions.

Dataset Description:

Examples:

1. Example 1:

- *Text*: "A subsection of retarded Hungarians? Ohh boy..."
- *Hate Speech Index*: [1]
- *Response*: "I don't see a reason why it's okay to insult..."

2. Example 2:

- *Text*: "> 'y'all hear sumn?' by all means I live in a small town..."
- *Hate Speech Index*: [3]
- *Response*: "Persons with disabilities is the accepted term..."

3. Example 3:

- *Text*: "Grammatical errors, overt racism, child prostitution..."
- *Hate Speech Index*: NaN
- *Response*: NaN

Benchmark Dataset

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5020 entries, 0 to 5019
Data columns (total 4 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   id                5020 non-null    object 
 1   text               5020 non-null    object 
 2   hate_speech_idx   3847 non-null    object 
 3   response          3847 non-null    object 
dtypes: object(4)
memory usage: 157.0+ KB
```

Data Visualization

- Dataset Sizes and Distributions:

Total Samples: 22,841

Class Distribution:

- *Training Count*

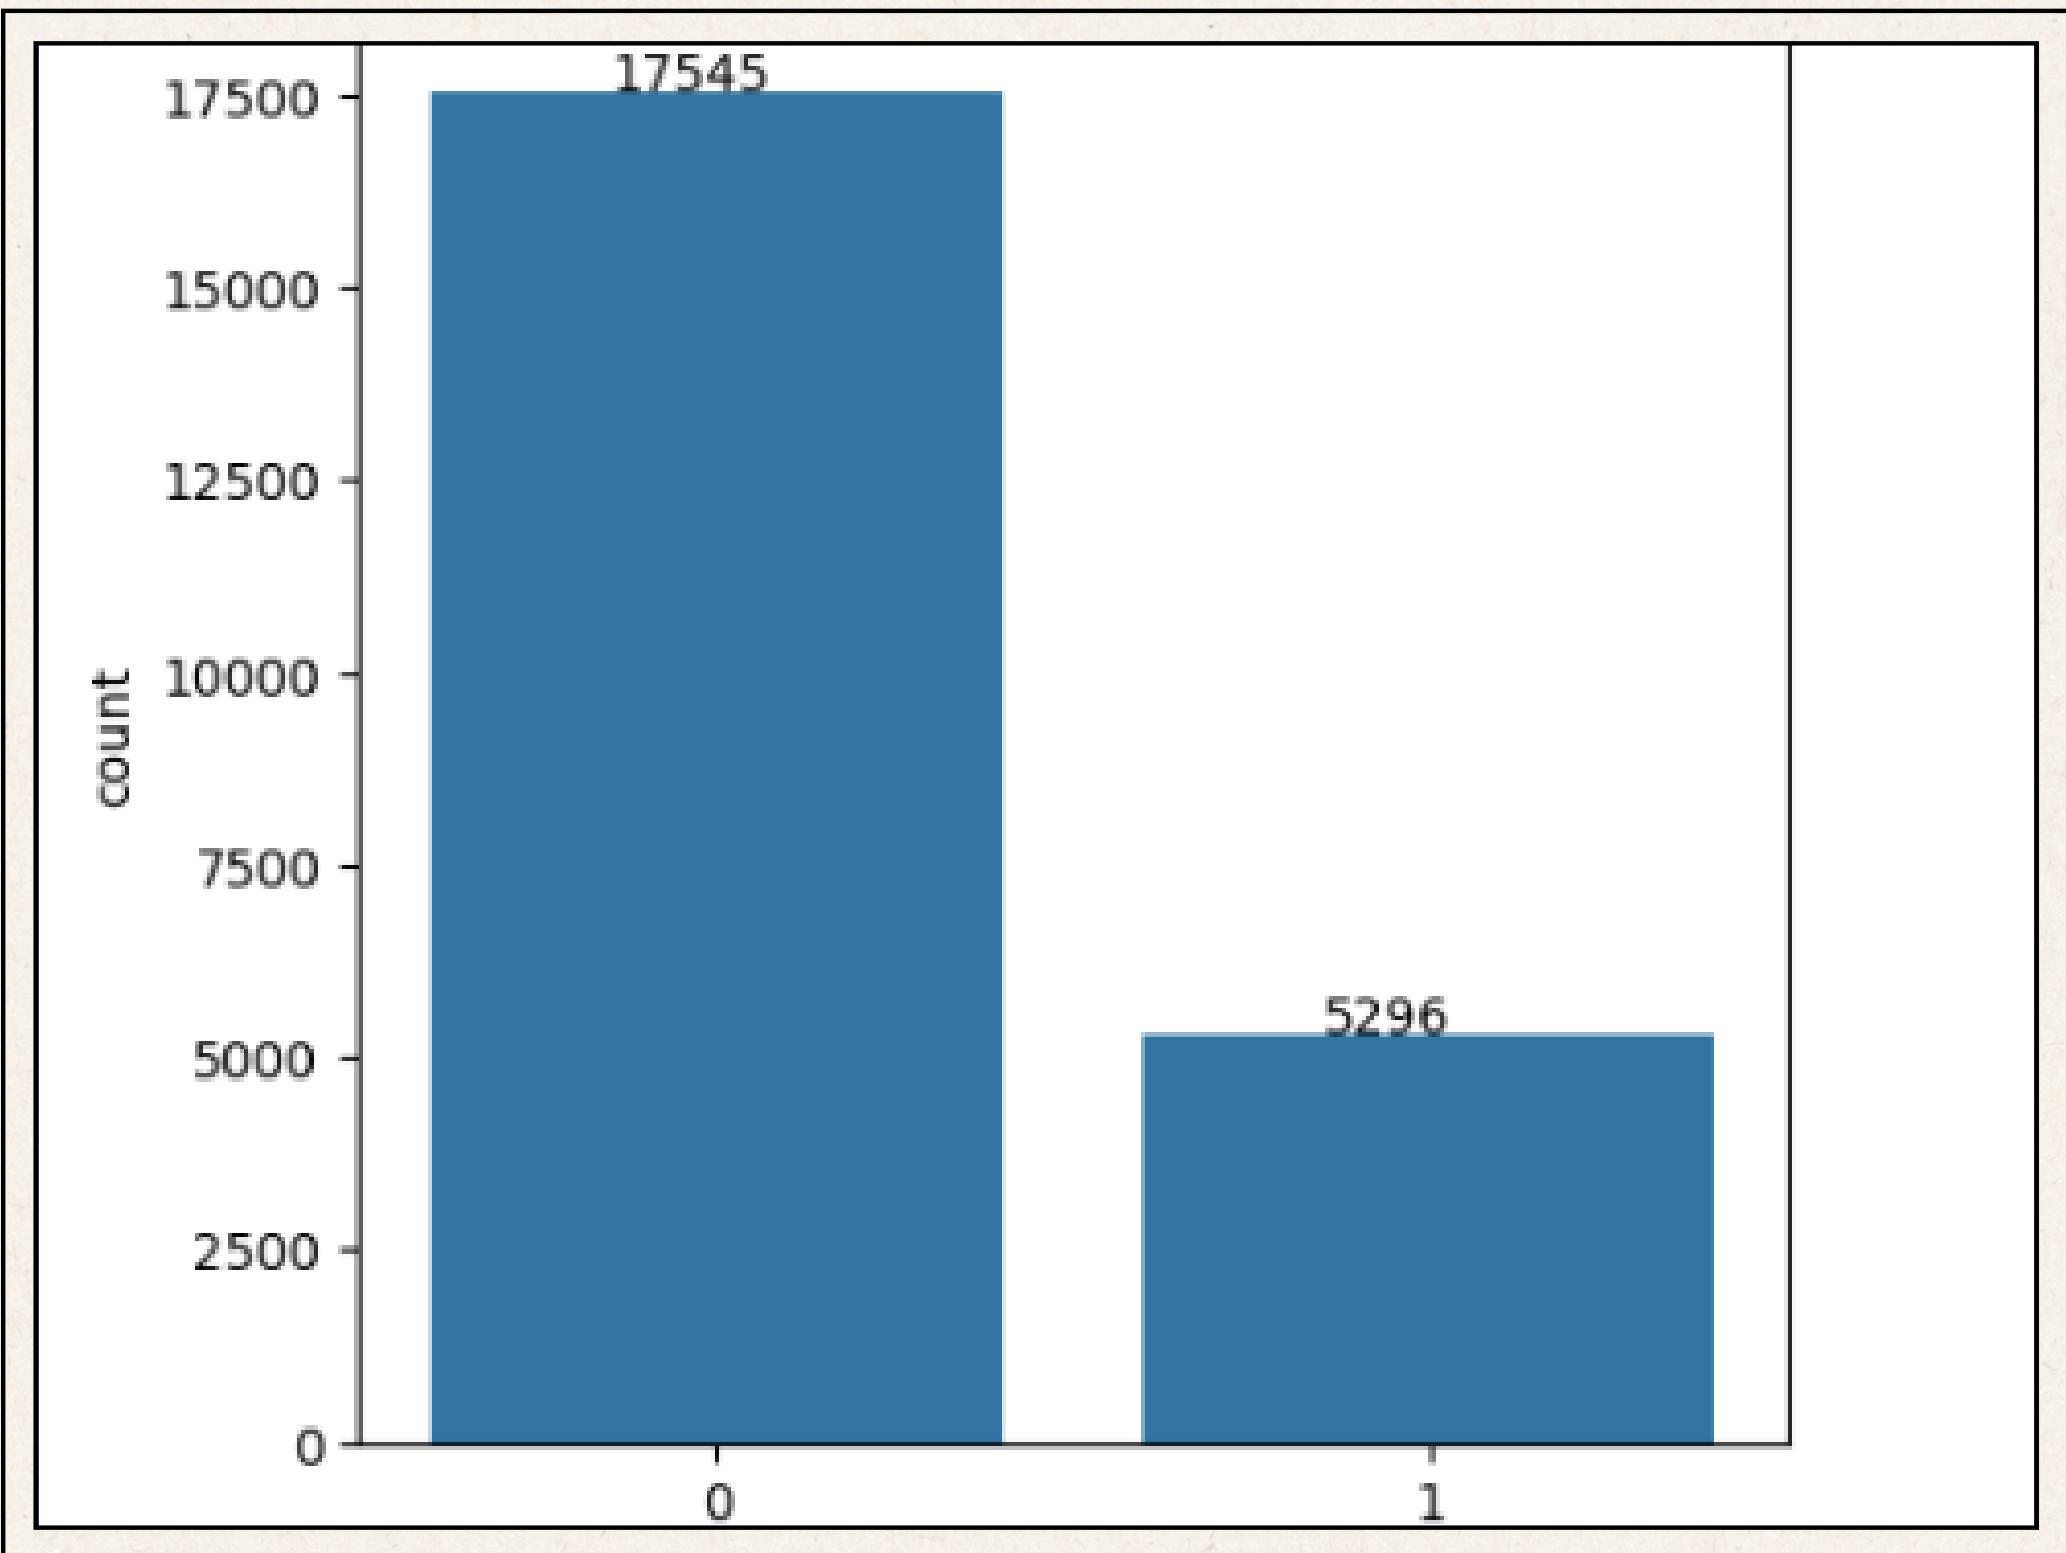
Non-Hate (0): 13,506 (76%)

Hate (1): 3,408 (77%)

- *Testing Count*

Non-Hate (0): 4,254 (24%)

Hate (1): 1,032 (23%)



Data Preprocessing

1. Importing Necessary Libraries:

- o Loaded essential libraries like pandas, numpy, sklearn, etc.
- o Imported the dataset into a dataframe for analysis.

2. Checking for Missing Values:

- o Identified and handled any missing values in the data.

3. Data Cleaning:

- o Removed numbers and chat words.
- o Handled emojis.
- o Applied autocorrect for text normalization.

4. Handling Imbalanced Data:

- o Applied SMOTE to balance the training data.

- o Class Distribution after SMOTE:

Non-Hate Speech (0): 13,539

Hate Speech (1): 13,539



Tokenization and Embedding Techniques

Tokenization Techniques:

1. Tokenization:

- Splitting text into individual words or tokens.

Embedding Techniques:

1. FastText (for ML Models):

- Usage: Applied FastText embeddings to represent words as dense vectors for machine learning models.

2. Word Embedding (for DL Models):

- Description: Converts words into continuous vector representations.
- Technique: Used pre-trained word embeddings or trained embeddings during model training.
- Advantages: Captures semantic meaning of words, improving model performance for text classification tasks.

Modeling

Various Machine Learning Models:

- Random Forest: Robust to overfitting, handles high-dimensional data, provides feature importance.
- Logistic Regression: Simple, interpretable, quick to train.
- Support Vector Machine (SVM): Effective for binary classification, robust with high-dimensional data.
- Gradient Boosting: Combines base estimators for improved accuracy, provides feature importance.
- K-Nearest Neighbors (KNN): Simple, effective for small datasets, predicts based on nearest data points.
- SGD Classifier: Efficient for large datasets.

Deep Learning Models:

- Convolutional Neural Network (CNN): Captures spatial and contextual information, performs well in text classification.
- Long Short-Term Memory (LSTM): Suitable for sequential data, captures long-term dependencies.

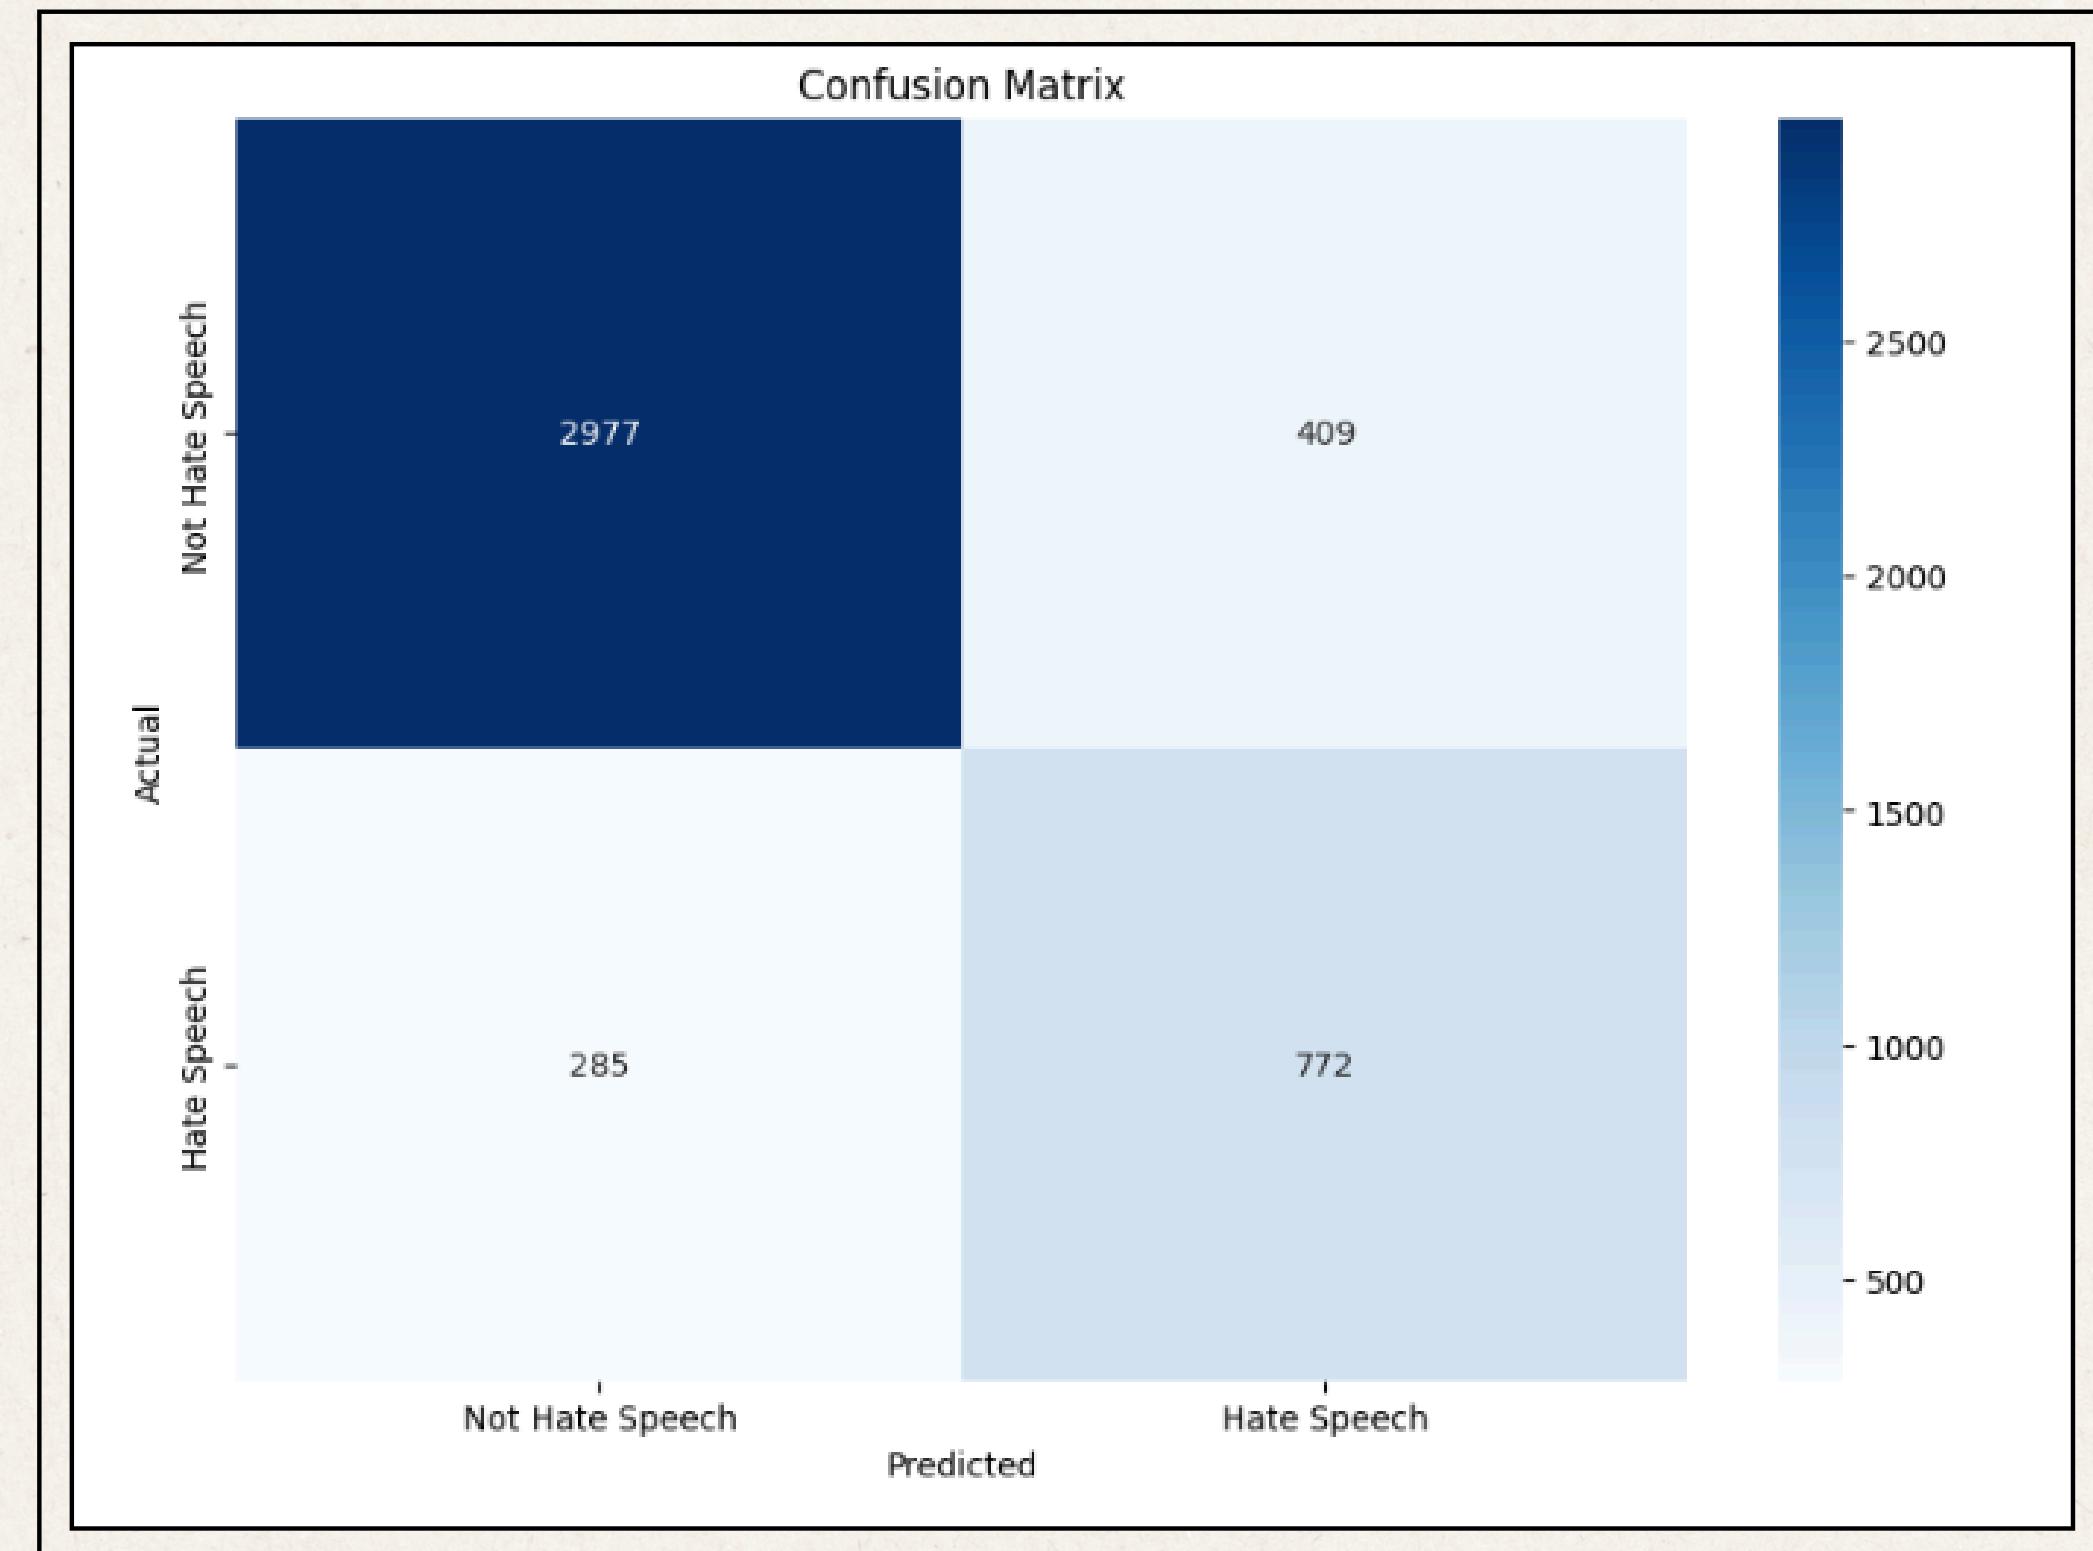
Evaluation Metrics

We used several key metrics to evaluate our models:

- **Accuracy:** Measured the overall correctness of the model.
- **F1 Score:** Balanced precision and recall, especially important for imbalanced datasets.
- **Confusion Matrix:** Provided a detailed breakdown of model performance across classes

DL Model (LSTM):

- Accuracy: 84.38%
- Precision: 65.37%
- Recall: 73.04%
- F1 Score: 68.99%



Thank You